

TU-Delft Deep Learning course 2018-2019

04.CNN

27 Feb 2019



Delft University of Technology

Lecturer: Jan van Gemert

Questions?

Logistics:

- Questions
- Presentations
- Lab

Questions?

Logistics:

- Questions
- Presentations
- Lab

Today: Convolutional Neural Networks

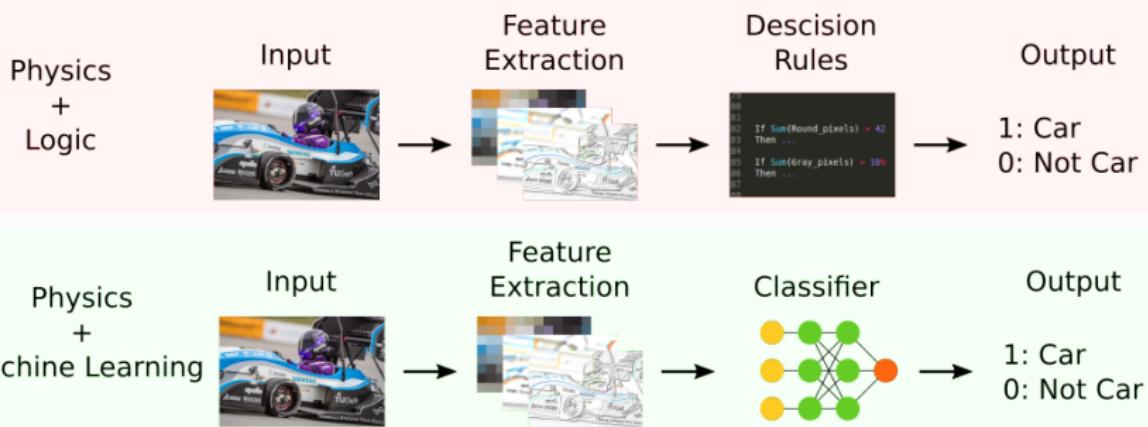
- Convolution
- Convolutional Network outline
- Relation to feed forward
- Properties of CNNs

Chapters: 9.1, 9.2, 9.3, 9.5

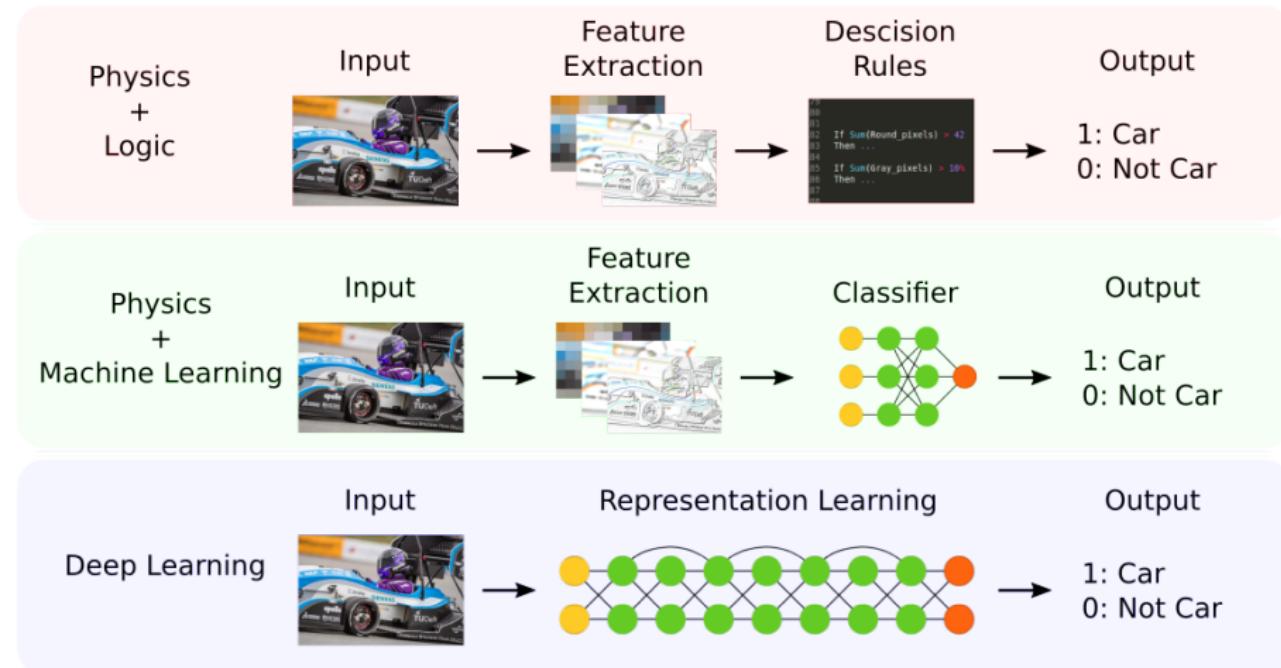
Feature extraction vs deep learning



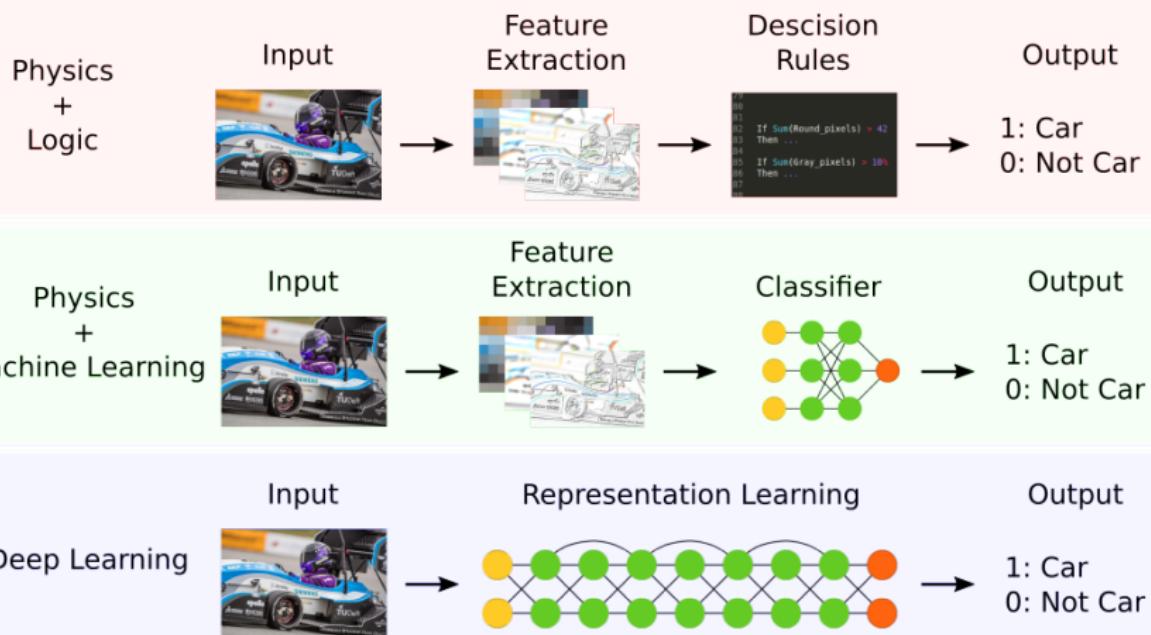
Feature extraction vs deep learning



Feature extraction vs deep learning



Feature extraction vs deep learning



End-to-End learning: End goal (output) used to learn feature extraction (input)

A bit of image processing

Q: How to get rid of noisy pixels?



A bit of image processing

Q: How to get rid of noisy pixels?



A: Simple solution: replace pixel by neighborhood average

Moving Neighborhood Average

$F(x, y)$

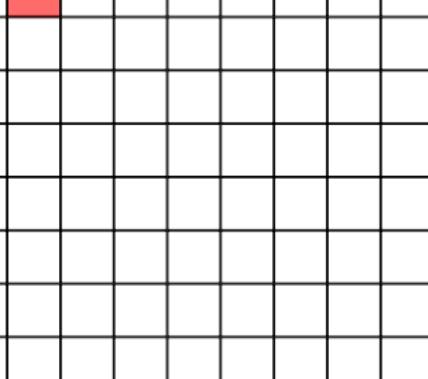
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	0	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

$G(x, y)$

Moving Neighborhood Average

$$F(x, y)$$

$$G(x, y)$$

A 10x10 grid of squares. The first column contains 10 white squares. The second column contains 9 white squares, with the top one being red. The remaining 8 columns each contain 10 white squares.

Moving Neighborhood Average

$F(x, y)$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

$G(x, y)$

			0								

Moving Neighborhood Average

$F(x, y)$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

$G(x, y)$

			0								

Moving Neighborhood Average

$F(x, y)$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

$G(x, y)$

			0	10							

Moving Neighborhood Average

$F(x, y)$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G(x, y)$

			0	10						

Moving Neighborhood Average

$F(x, y)$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G(x, y)$

	0	10	20							

Moving Neighborhood Average

$F(x, y)$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G(x, y)$

	0	10	20						

Moving Neighborhood Average

$F(x, y)$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G(x, y)$

	0	10	20	30					

Moving Neighborhood Average

$F(x, y)$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	0	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

$G(x, y)$

	0	10	20	30	30	30	20	10			
	0	20	40	60	60	60	40	20			
	0	30	60	90	90	90	60	30			
	0	30	50	80	80	90	60	30			
	0	30	50	80	80	90	60	30			
	0	20	30	50	50	60	40	20			
	10	20	30	30	30	30	20	10			
	10	10	10	0	0	0	0	0			

Moving Neighborhood Average

Q: What do you notice?

$F(x, y)$

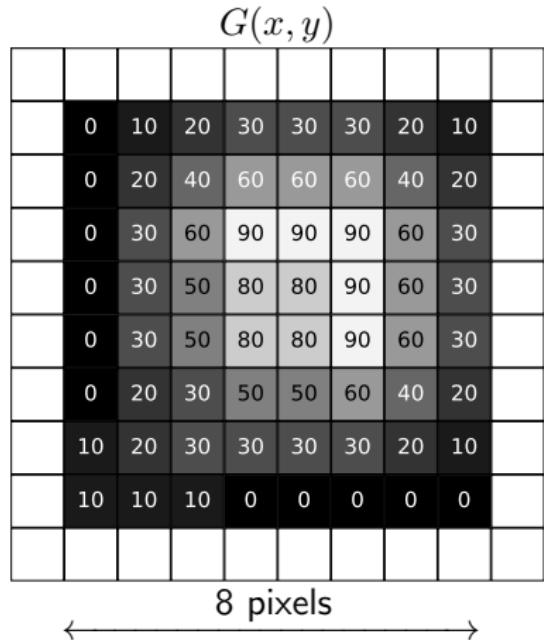
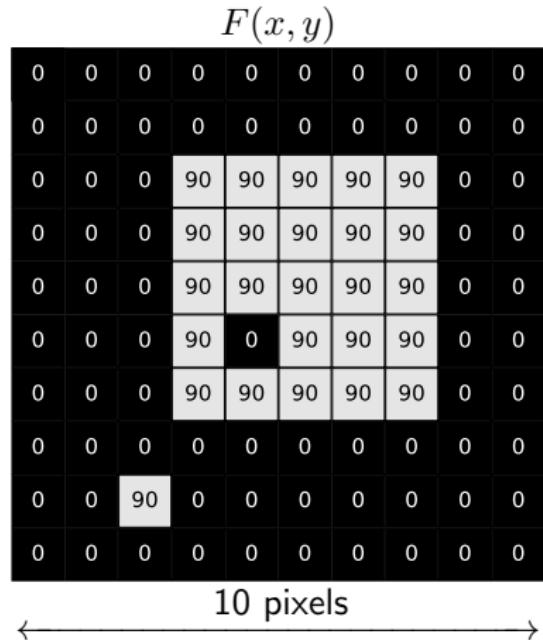
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	0	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

$G(x, y)$

	0	10	20	30	30	30	20	10			
	0	20	40	60	60	60	40	20			
	0	30	60	90	90	90	60	30			
	0	30	50	80	80	90	60	30			
	0	30	50	80	80	90	60	30			
	0	20	30	50	50	60	40	20			
	10	20	30	30	30	30	20	10			
	10	10	10	0	0	0	0	0			

Moving Neighborhood Average

Q: What do you notice?



A: 2 pixels lost to boundary (1 on each side)

Convolution

Chapter 9.1

Generalize to a kernel: **Multiply weights variables per pixel and sum**

$$F(x, y) \star H(x, y) = G(x, y)$$

The diagram illustrates the convolution operation between a feature map $F(x, y)$ and a kernel $H(x, y)$ to produce a result $G(x, y)$.

Feature Map $F(x, y)$:

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	90	0	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

Kernel $H(x, y)$:

$$\begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}$$

Result $G(x, y)$:

			0	10	20	30		

Q: What values to fill in kernel H to obtain a moving neighborhood average?

Convolution

Chapter 9.1

Generalize to a kernel: **Multiply weights variables per pixel and sum**

$$F(x, y) \star H(x, y) = G(x, y)$$

Input image $F(x, y)$ (3x3 kernel):

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	90	0	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

Kernel $H(x, y)$:

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Output image $G(x, y)$:

			0	10	20	30		

Q: What values to fill in kernel H to obtain a moving neighborhood average?

Practice with kernels

 $F(x, y)$  \star $H(x, y)$ $=$ $G(x, y)$

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

Practice with kernels

 $F(x, y)$  \star
 $H(x, y)$

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

 $=$
 $G(x, y)$ 

Practice with kernels

 $F(x, y)$  $\star \quad H(x, y) \quad =$ $G(x, y)$

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

Practice with kernels

 $F(x, y)$  \star
 $H(x, y)$

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

 $=$
 $G(x, y)$ 

Practice with kernels

 $F(x, y)$  $\star \quad H(x, y) =$ $G(x, y)$

$$\begin{bmatrix} 0 & 0 & 0 \\ -1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

Practice with kernels

 $F(x, y)$  \star
 $H(x, y)$

$$\begin{bmatrix} 0 & 0 & 0 \\ -1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

 $=$
 $G(x, y)$ 

Practice with kernels

 $F(x, y)$  $\star \quad H(x, y) \quad =$ $G(x, y)$

$$\begin{bmatrix} 0 & 0 & 0 \\ 1 & 0 & -1 \\ 0 & 0 & 0 \end{bmatrix}$$

Practice with kernels

 $F(x, y)$  \star
 $H(x, y)$

$$\begin{bmatrix} 0 & 0 & 0 \\ 1 & 0 & -1 \\ 0 & 0 & 0 \end{bmatrix}$$

 $=$
 $G(x, y)$ 

Practice with kernels

(Red = positive; blue is negative)

$$F(x, y)$$



*

$$H(x, y)$$



=

$$G(x, y)$$

Practice with kernels

(Red = positive; blue is negative)

$$F(x, y)$$



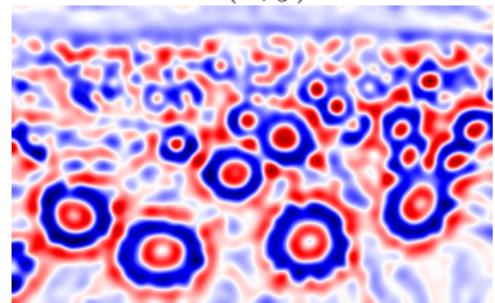
*

$$H(x, y)$$



=

$$G(x, y)$$



Practice with kernels

(Red = positive; blue is negative)

$$F(x, y)$$



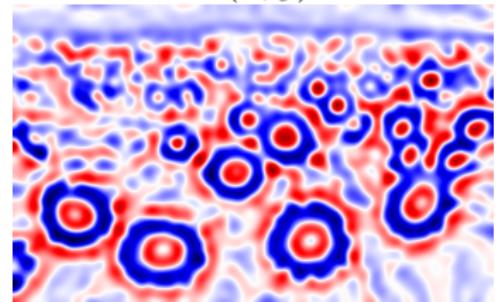
*

$$H(x, y)$$



=

$$G(x, y)$$



Q: How about a smaller kernel?



Practice with kernels

(Red = positive; blue is negative)

$$F(x, y)$$



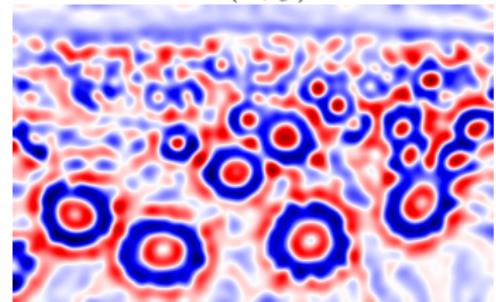
*

$$H(x, y)$$

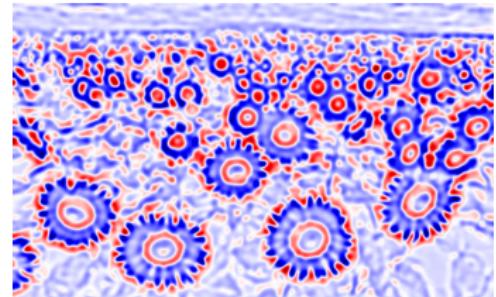


=

$$G(x, y)$$



Q: How about a smaller kernel?



Practice with kernels

(Red = positive; blue is negative)

$$F(x, y)$$



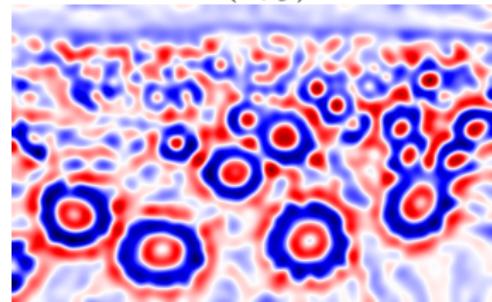
*

$$H(x, y)$$

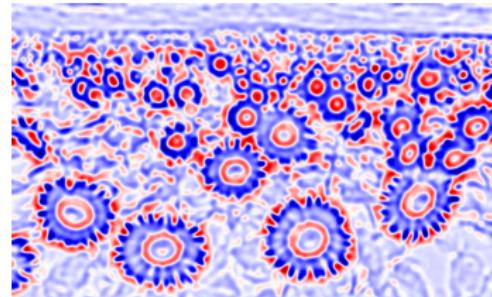


=

$$G(x, y)$$



Q: How about a smaller kernel?



Q: How about a smaller input image?

Practice with kernels

(Red = positive; blue is negative)

$$F(x, y)$$



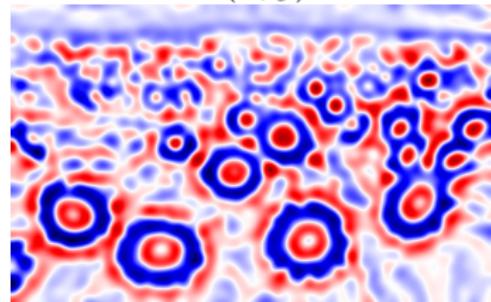
*

$$H(x, y)$$

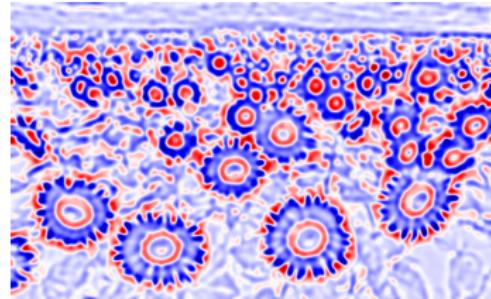


=

$$G(x, y)$$



Q: How about a smaller kernel?



Q: How about a smaller input image?

A: Smaller image = larger kernel.

Where is Waldo?



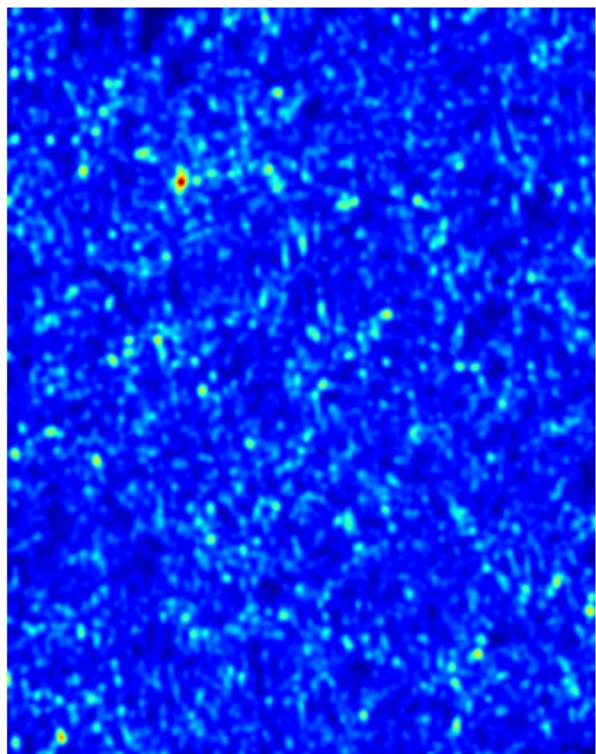
Where is Waldo?



Use this normalized kernel:



Where is Waldo?



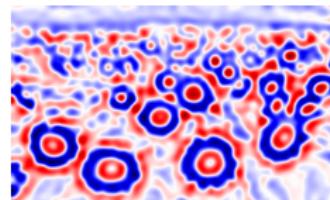
Representation learning



*



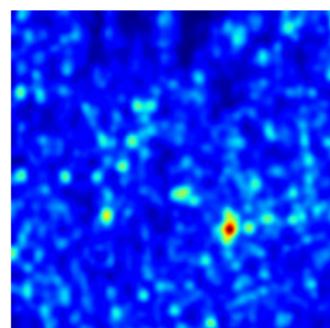
=



*



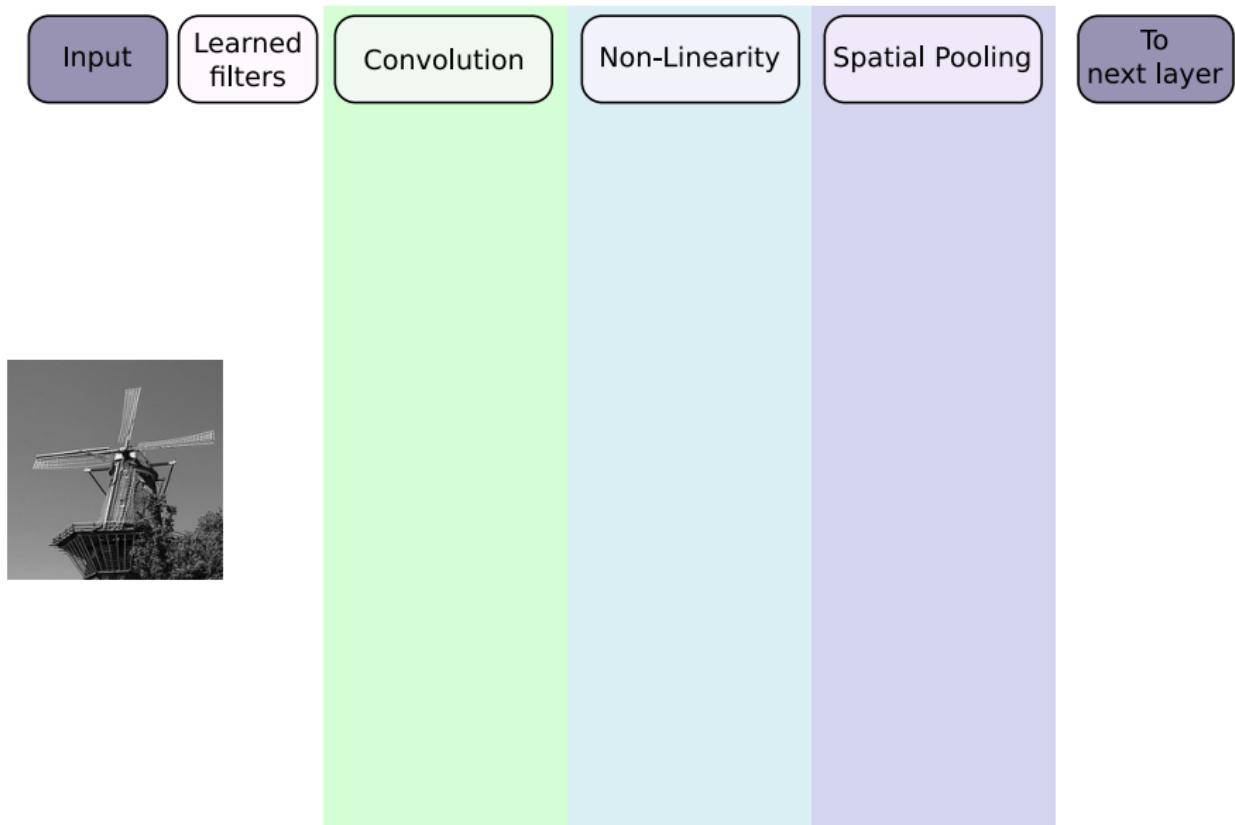
=



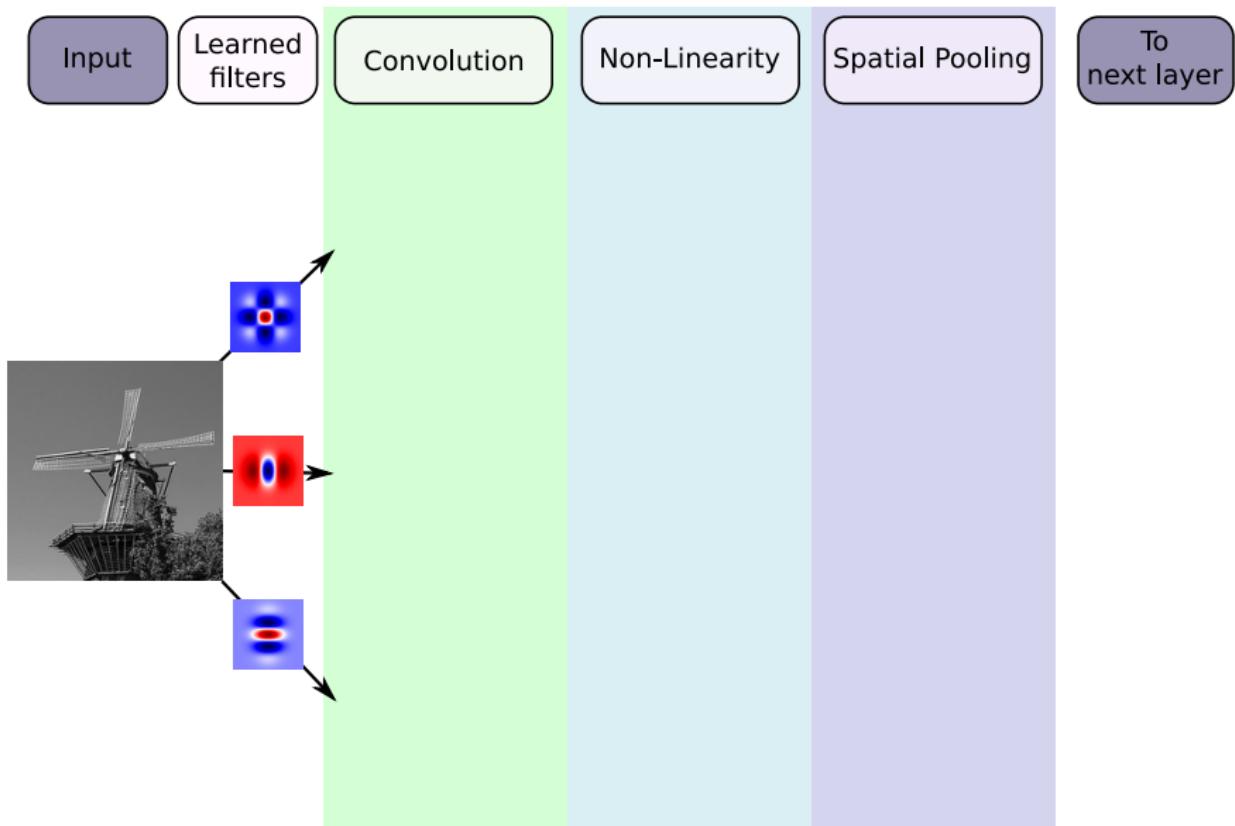
- Kernel weights are feature detectors
- Learning weights = Learning features
- Convnet learns the feature representation

Questions?

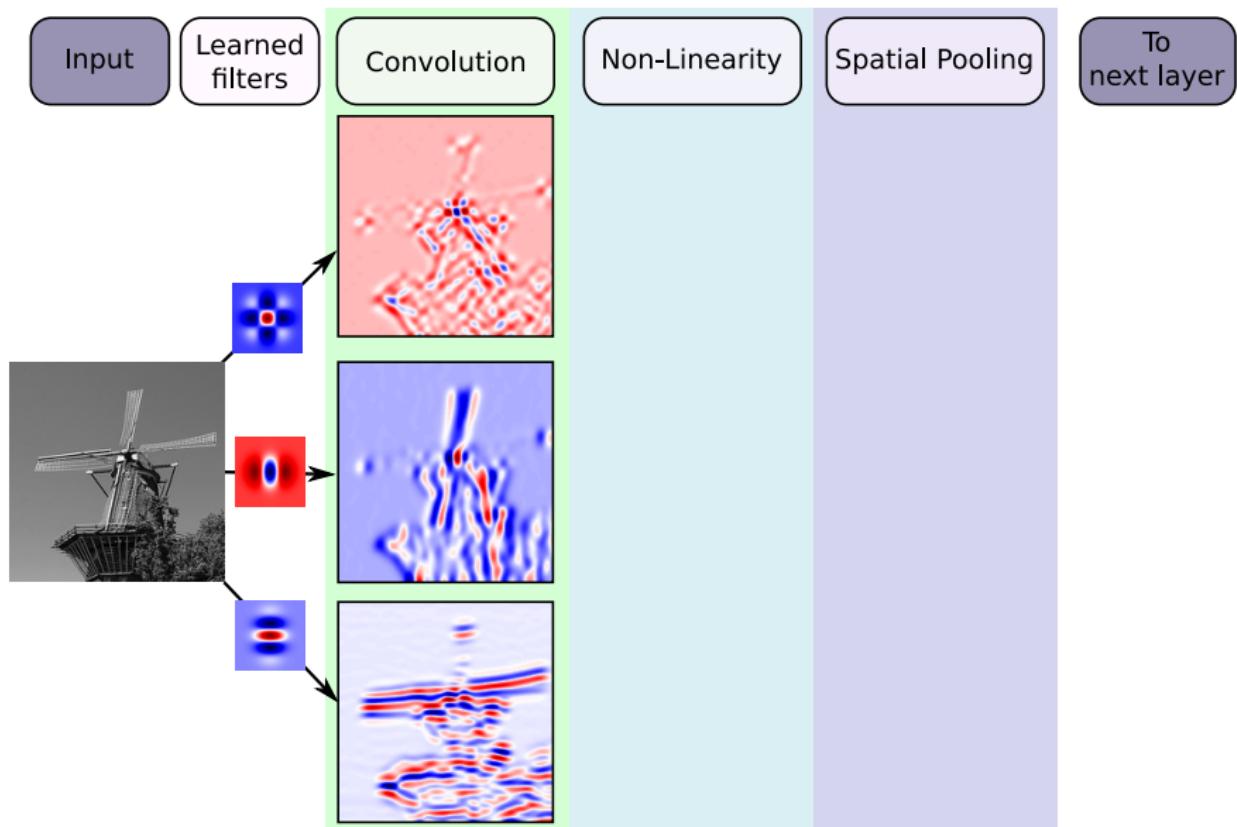
Convolutional Network



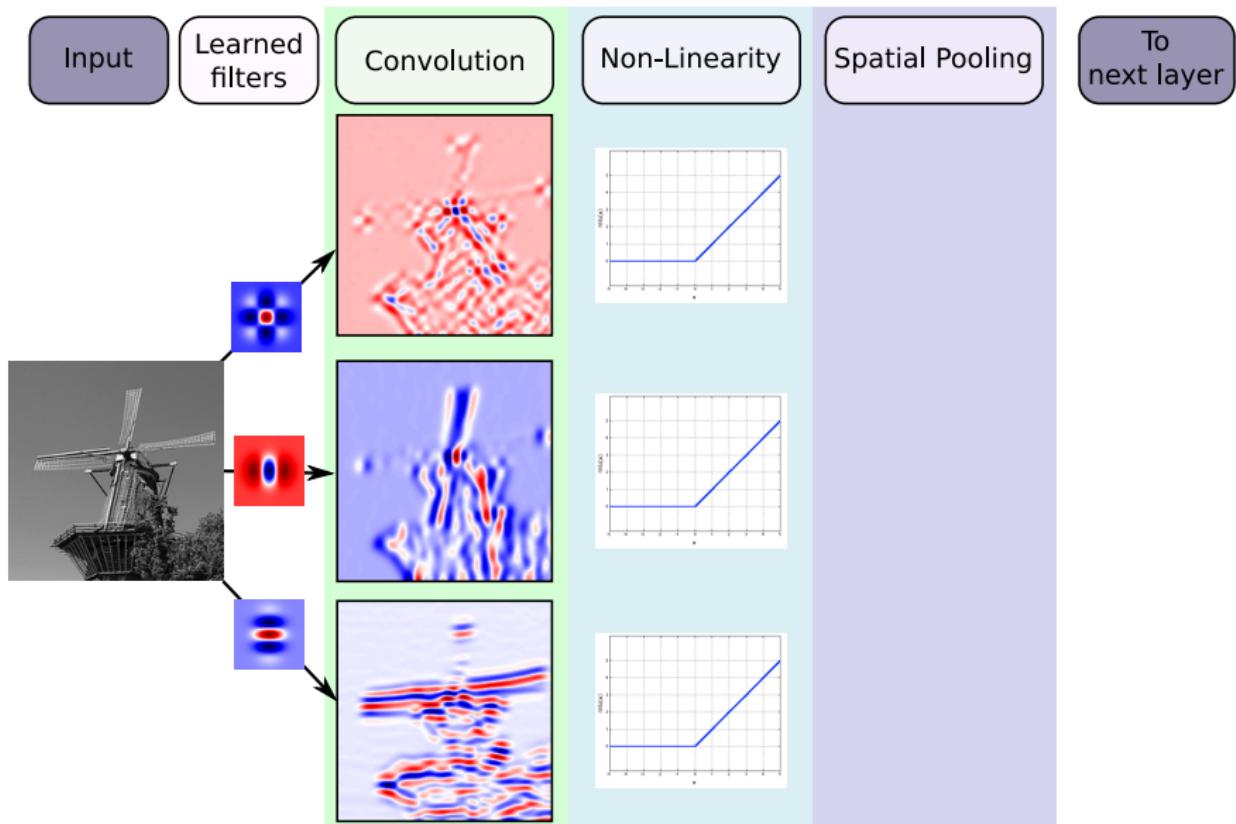
Convolutional Network: Filter



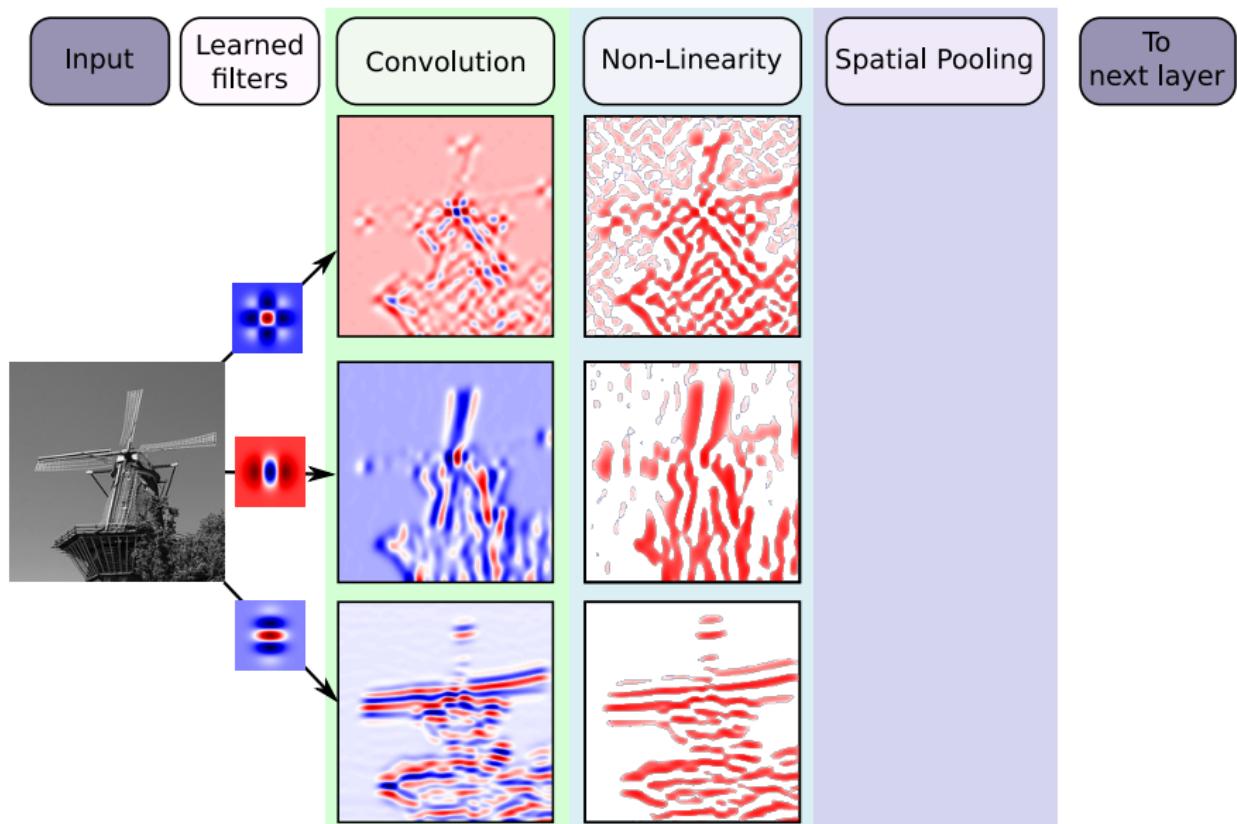
Convolutional Network: Featuremaps



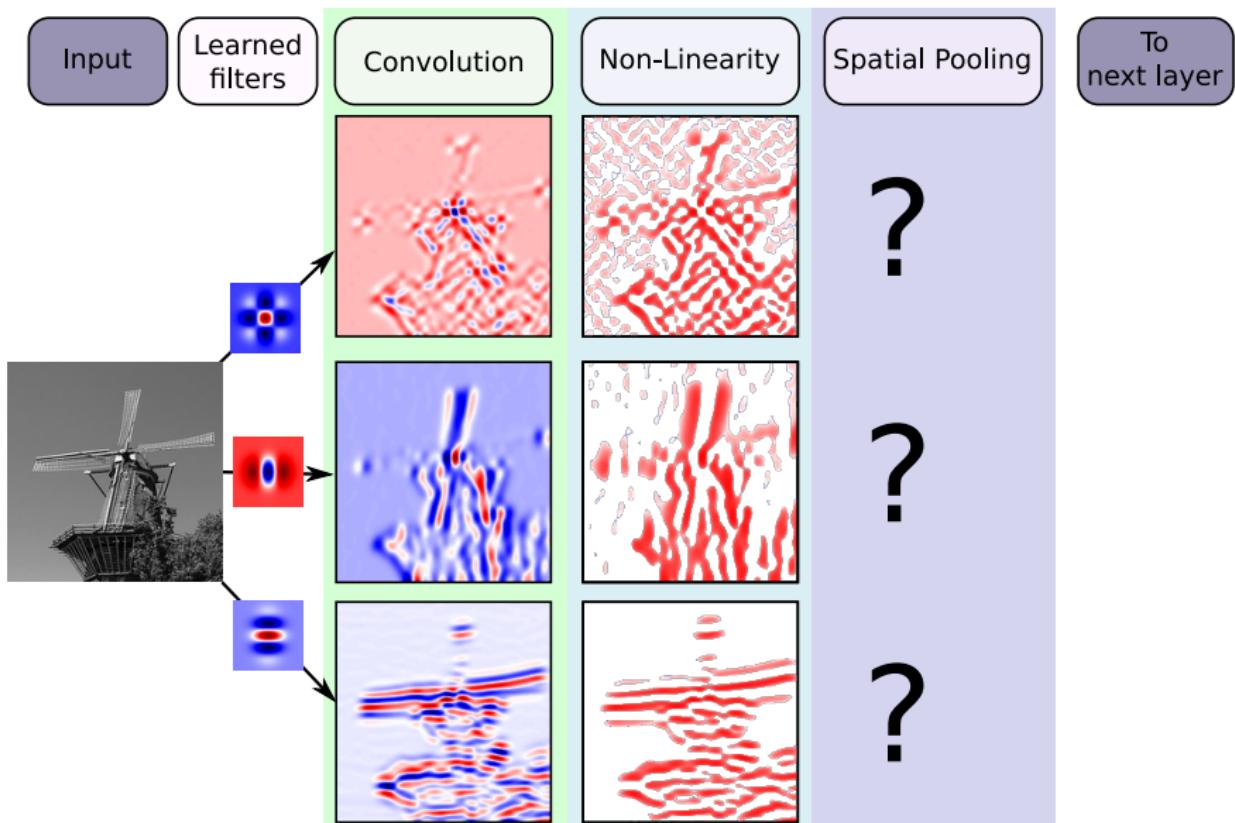
Convolutional Network: ReLu ($f(x) = \max(0, x)$)



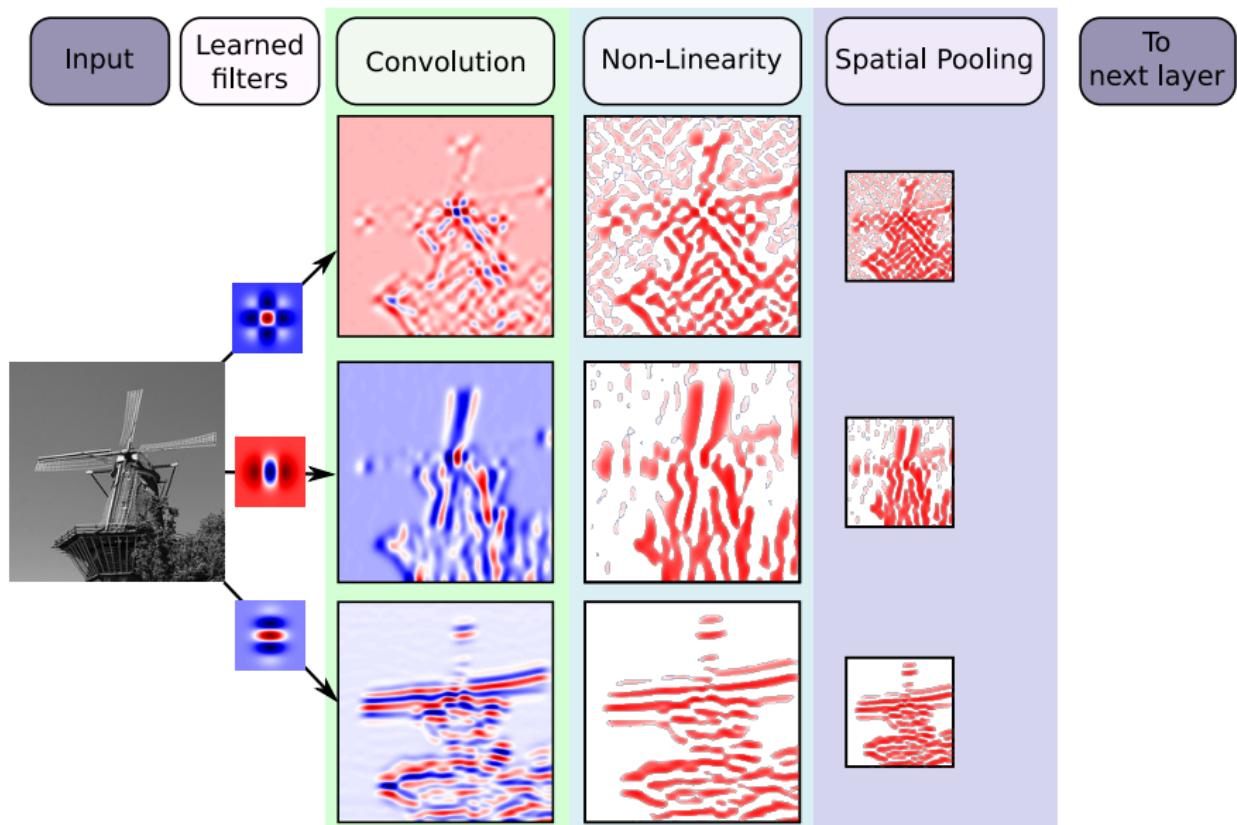
Convolutional Network: All negative values removed



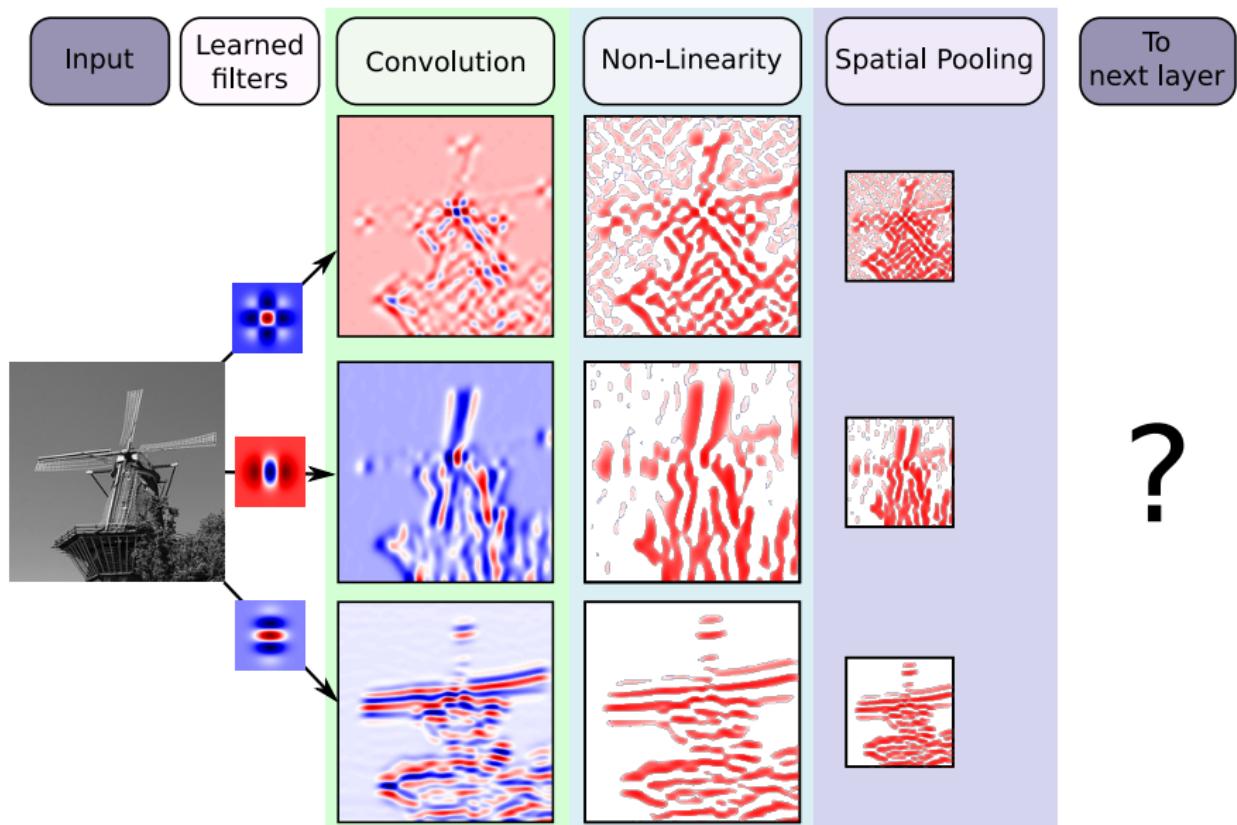
Convolutional Network: 2×2 max, 2×2 sub-sample



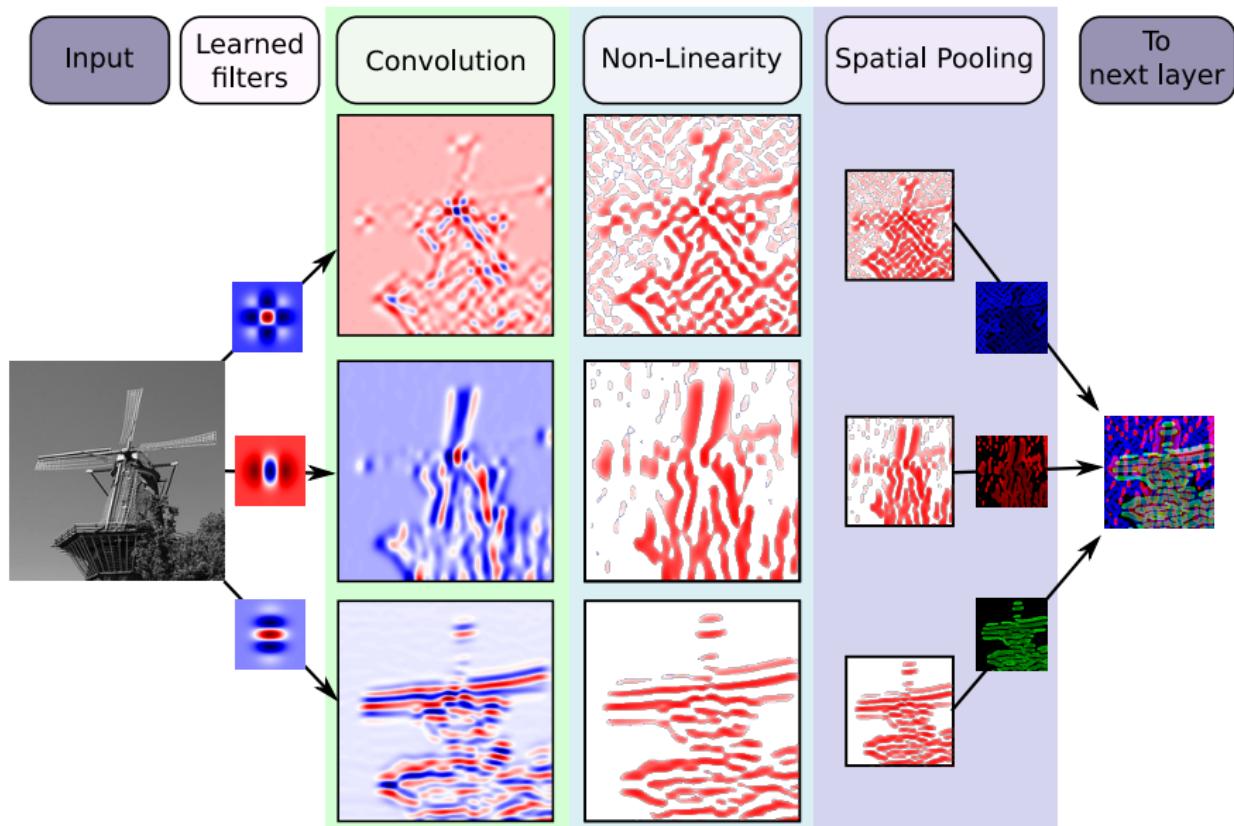
Convolutional Network: Smaller feature maps



Convolutional Network: To the next layer

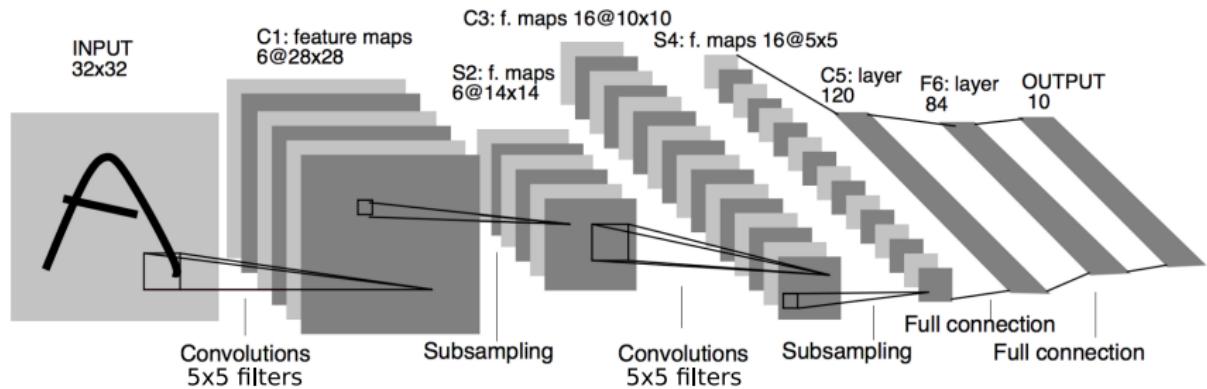


Convolutional Network: Featuremaps as new channels



Questions?

Questions?



Q: Can you explain this CNN?